

Conversational AI

Twisha Shah

Georgia Institute of Technology

tshah93@gatech.edu

Abstract

Conversational AI has revolutionized human-machine interaction, enabling natural language dialogue between humans and computers. This survey provides a comprehensive overview of the advancements in Conversational AI, from traditional rule-based systems to modern end-to-end neural network architectures. We discuss the key challenges, evaluation metrics, and applications of Conversational AI, highlighting its significance in transforming customer service, healthcare, and education. Our analysis reveals the future scope of Conversational AI, emphasizing the need for multimodal interaction, emotional intelligence, and Explainability. This survey serves as a comprehensive resource for researchers, developers, and practitioners seeking to understand the complexities and possibilities of Conversational AI.

1 Keywords

Conversational AI, Natural Language Processing, Dialogue Systems, Chatbots, Virtual Assistants, Reinforcement Learning, Transformers, BERT, Dialogue Management, Intent Detection, Sentiment Analysis, Emotional Intelligence, Multimodal Interaction.

2 Introduction

Conversational Artificial Intelligence (CAI) has emerged as a paradigm-shifting technology, revolutionizing human-machine interaction by enabling machines to engage in natural language dialogue with humans. This burgeoning field has witnessed exponential growth, driven by breakthroughs in Natural Language Processing (NLP), machine learning, and cognitive computing. The

convergence of these technologies has given rise to sophisticated conversational systems, transforming various domains, including customer service, healthcare, education, and human-computer interaction. The increasing demand for human-like interaction with machines has been a primary driver of CAI's development. Traditional interface paradigms, reliant on keyboards, touchscreens, and graphical user interfaces (GUIs), have limitations in capturing nuanced user intent, emotions, and context. CAI addresses these limitations by facilitating user interaction in natural language, either through voice or text, thereby enabling more intuitive and effective human-machine collaboration. The scientific community has witnessed a significant surge in CAI research, with a focus on developing more sophisticated dialogue management systems, incorporating cognitive architectures, and integrating multimodal interaction modalities. However, the complexity and diversity of CAI systems have created a need for a comprehensive overview of the field, highlighting its historical evolution, theoretical foundations, and practical applications.

This survey aims to provide a thorough and systematic examination of CAI, bridging the gap between researchers, practitioners, and stakeholders. We strive to:

- Provide a historical context for CAI, tracing its development and key milestones.
- Offer a detailed analysis of CAI architectures, including rule-based, statistical, and deep learning models.
- Discuss evaluation metrics and methods for assessing CAI performance.
- Showcase real-world applications and case studies of CAI.
- Identify current challenges, limitations, and open research questions.

- Explore future directions, emerging trends, and potential breakthroughs.

By presenting a comprehensive and structured overview of CAI, this survey seeks to facilitate a deeper understanding of the field's current state, its potential, and its impact on various domains.

3 Background

The concept of Conversational Artificial Intelligence (CAI) has its roots in the early 20th century, with the pioneering work of Alan Turing. Turing's 1950 paper, "Computing Machinery and Intelligence," proposed the Turing Test, a measure of a machine's ability to exhibit intelligent behavior equivalent to, or indistinguishable from, that of a human. This idea laid the foundation for the development of chatbots and conversational systems. In the 1960s, the first chatbot, ELIZA, was developed by Joseph Weizenbaum. ELIZA's ability to simulate a conversation by using a set of pre-defined rules and responses marked the beginning of CAI research. The 1970s and 1980s saw the emergence of rule-based systems, which relied on hand-crafted knowledge graphs and logical rules to generate responses. These systems were limited by their inability to learn from data and adapt to new situations. The advent of machine learning in the 1990s and 2000s enabled the development of statistical models, such as Hidden Markov Models (HMMs) and Conditional Random Fields (CRFs). These models improved CAI's ability to recognize patterns in language and generate more coherent responses. The recent surge in deep learning techniques, particularly Recurrent Neural Networks (RNNs) and Transformers, has revolutionized CAI. These models have enabled the development of more sophisticated dialogue management systems, capable of learning from large datasets and generating human-like responses. Key milestones in CAI's evolution include:

- 1950: Alan Turing proposes the Turing Test
- 1966: Joseph Weizenbaum develops ELIZA, the first chatbot
- 1970s: Rule-based systems emerge
- 1990s: Statistical models, such as HMMs and CRFs, are developed
- 2000s: Machine learning and deep learning techniques become prominent in CAI research

- 2018: BERT is introduced, marking a significant improvement in language understanding
- 2020: Conversational Transformers and explainability techniques gain attention

Recent updates in CAI include:

- Pre-trained Language Models: The development of pre-trained language models like BERT (2018), RoBERTa (2019), and XLNet (2019) has significantly improved CAI's language understanding capabilities.
- Conversational Transformers: The introduction of conversational Transformers, such as Dialogue Transformers (2020) and Conversational BERT (2020), has enabled more efficient and effective dialogue management.
- Explainability and Transparency: Research on explainability and transparency in CAI has gained momentum, with techniques like attention visualizations and model interpretability methods being developed.
- Multimodal Interaction: The integration of multimodal interaction modalities, such as vision, speech, and gesture recognition, has expanded CAI's applications in areas like human-robot interaction and virtual assistants.
- Edge AI and Real-time Processing: The increasing availability of edge AI computing resources and real-time processing capabilities has enabled CAI systems to be deployed in resource-constrained environments, such as smart home devices and autonomous vehicles.

4 Architecture and Approaches

Conversational AI systems typically consist of several components, including: Natural Language Processing (NLP): responsible for text or speech input processing, including tokenization, part-of-speech tagging, named entity recognition, and sentiment analysis. Dialogue Management: determines the system's response based on the user's input and the conversation context. Knowledge Retrieval: retrieves relevant information from a knowledge base or database to inform the system's response. Response Generation: generates

the system's response, either by retrieving a pre-existing response or generating a new one using machine learning algorithms.

Conversational AI architectures rely on *dialogue management systems* to govern the flow of conversation. These systems often employ finite state machines, which use states, transitions, and actions to manage dialogue. Additionally, frame-based systems utilize a structured representation of knowledge to manage dialogue, while plan-based systems leverage planning algorithms to generate dialogue responses. Hybrid approaches combine multiple techniques to manage dialogue, offering flexibility and robustness.

Natural Language Processing (NLP) pipelines play a crucial role in conversational AI, enabling systems to understand and interpret user input. Text preprocessing is a critical step, involving techniques like tokenization, stemming, and lemmatization to normalize text data. Part-of-speech tagging identifies word types, while named entity recognition extracts specific entities like names and locations. Dependency parsing analyzes sentence structure, providing context for accurate understanding.

Machine learning algorithms are integral to conversational AI, enabling systems to learn and improve over time. Supervised learning algorithms, such as support vector machines and random forests, are commonly used, while deep learning algorithms like recurrent neural networks, long short-term memory networks, and transformers offer advanced capabilities. Reinforcement learning algorithms, which learn from trial and error, are also applied in conversational AI to optimize system performance and user experience.

4.1 Rule-Based Systems

Rule-Based Systems rely on hand-crafted rules and knowledge graphs to manage conversation flow. These systems use a set of pre-defined rules to match user input and generate responses.

The rule-based approach can be represented as:

$$R = f(I, K)$$

where R is the response, I is the input, K is the knowledge base, and f is the rule-based function.

4.1.1 Architecture

- Knowledge Graph: A database of pre-defined knowledge and rules.

- Rule Engine: A module that applies rules to generate responses.
- Response Generator: A module that generates responses based on the output of the rule engine.

4.1.2 Advantages

- Easy to implement and maintain
- Can handle simple conversations with well-defined rules
- Fast response generation

4.1.3 Limitations

- Limited in handling complex conversations with ambiguous or unclear input
- Requires manual rule creation and updating
- May not adapt to user behavior or context

4.2 Statistical Machine Translation (SMT)

SMT uses statistical models to translate user input into a response. These models are trained on large datasets of input-response pairs.

The SMT approach can be represented as:

$$P(R|I) = \operatorname{argmax} P(R, I)$$

where $P(R|I)$ is the probability of generating a response given the input, and $P(R, I)$ is the joint probability of the response and input.

4.2.1 Architecture

- Data Preprocessing: A module that preprocesses input data.
- Model Training: A module that trains the statistical model on the preprocessed data.
- Model Inference: A module that uses the trained model to generate responses.

4.2.2 Advantages

- Can handle complex conversations with nuanced input
- Can learn from large datasets
- Supports multiple languages

4.2.3 Limitations

- Requires large amounts of training data
- May not adapt to user behavior or context
- Can be computationally expensive

4.3 Deep Learning-based Systems

Deep Learning-based Systems use neural networks to learn from large datasets and generate responses. These networks can be trained on various tasks, such as language modeling, sentiment analysis, and dialogue generation.

The deep learning approach can be represented as:

$$R = f(I, \theta)$$

where R is the response, I is the input, θ is the model parameters, and f is the neural network function.

4.3.1 Architecture

- Embedding Layer: A module that converts input text into numerical representations.
- Encoder: A module that encodes the input text into a continuous representation.
- Decoder: A module that generates responses based on the encoded representation.
- Output Layer: A module that generates the final response.

4.3.2 Advantages

- Can learn from large datasets and adapt to user behavior
- Can handle complex conversations with nuanced input
- Supports multiple tasks and languages

4.3.3 Limitations

- Requires significant computational resources
- May require large amounts of training data
- Can be challenging to interpret and debug

4.4 Sequence-to-Sequence (Seq2Seq) Models

Seq2Seq models use an encoder-decoder framework to generate responses. The encoder processes the input sequence, and the decoder generates the response sequence.

The Seq2Seq approach can be represented as:

$$R = f(E(I), D)$$

where R is the response, I is the input, E is the encoder, D is the decoder, and f is the Seq2Seq function.

4.4.1 Advantages

- Can handle sequential data, such as text or speech
- Can learn from large datasets
- Supports multiple tasks and languages

4.4.2 Limitations

- Can be computationally expensive
- May require large amounts of training data
- Can suffer from overfitting or underfitting

4.5 Attention-based Models

Attention-based Models use attention mechanisms to focus on specific parts of the input when generating a response. These mechanisms can be used in conjunction with Seq2Seq models.

The attention-based approach can be represented as:

$$R = f(E(I), A)$$

where R is the response, I is the input, E is the encoder, A is the attention mechanism, and f is the attention-based function.

4.5.1 Advantages

- Can handle complex conversations with nuanced input
- Can learn from large datasets
- Supports multiple tasks and languages

4.5.2 Limitations

- Can be computationally expensive
- May require large amounts of training data
- Can suffer from overfitting or underfitting

4.6 Transformer-based Models

Transformer-based Models use self-attention mechanisms to process input sequences in parallel. These models have achieved state-of-the-art results in various NLP tasks.

The transformer-based approach can be represented as:

$$R = f(E(I), S)$$

where R is the response, I is the input, E is the encoder, S is the self-attention mechanism, and f is the transformer-based function.

4.6.1 Advantages

- Can handle complex conversations with nuanced input
- Can learn from large datasets
- Supports multiple tasks and languages
- Parallelizable, reducing computational cost

4.6.2 Limitations

- Can be computationally expensive for large input sequences
- May require large amounts of training data
- Can suffer from overfitting or underfitting

4.7 Reinforcement Learning-based Models

Reinforcement Learning-based Models use reinforcement learning to optimize the system's response based on user feedback.

Reinforcement Learning (RL) is a crucial aspect of conversational AI, enabling dialogue systems to learn from user interactions and optimize their responses. In RL, the dialogue system acts as an agent, interacting with the environment (the conversation) to maximize a reward signal. The agent's goal is to learn a policy that maps states to actions, selecting responses that lead to the highest reward. The RL process involves user input, state updates, action selection, reward calculation, and policy updates. Algorithms like Q-Learning, SARSA (State-Action-Reward-State-Action), and Deep Q-Networks (DQN) are used to update the policy and value function, enabling the dialogue system to adapt to changing user behavior and preferences. RL brings significant benefits to conversational AI, including adaptability, optimization, and personalization. By learning from user feedback, RL enables dialogue systems to optimize conversation flow and response selection, leading to improved user experience and satisfaction. However, RL in conversational AI also poses challenges, such as defining effective reward signals, balancing exploration and exploitation, and handling conversational complexity. Despite these challenges, RL has the potential to revolutionize conversational AI, enabling more sophisticated and user-centric dialogue systems. These models can learn from interactions with users.

The reinforcement learning-based approach can be represented as:

$$R = f(I, U)$$

Approach	Complexity	Scalability
Rule-based	Low	High
Statistical	Medium	Medium
DL based	High	Low
Hybrid	High	Medium

Table 1: Different approaches and their complexity and scalability

where R is the response, I is the input, U is the user feedback, and f is the reinforcement learning-based function.

4.7.1 Advantages

- Can learn from user feedback and adapt to user behavior
- Can handle complex conversations with nuanced input
- Supports multiple tasks and languages

4.7.2 Limitations

- Requires user feedback, which can be challenging to obtain
- Can be computationally expensive
- May suffer from exploration-exploitation trade-offs

5 Evaluation Metrics

Evaluating Conversational AI systems requires a combination of automatic and human evaluation metrics. Automatic metrics provide quantitative measures, while human evaluation provides qualitative insights.

5.1 Automatic Metrics

5.1.1 Perplexity

Perplexity measures the system's ability to predict user input. A lower perplexity indicates better performance. Mathematically, perplexity is represented as:

$$P = 2^{(-\sum(p(x) \cdot \log_2(p(x)))}$$

where P is the perplexity and p(x) is the probability of the predicted response.

Example: Suppose we have a CAI model that predicts a response with a probability of 0.8. The perplexity would be:

$$PP = 2^{(-0.8 \cdot \log_2(0.8))} \approx 1.25$$

Interpretation: A perplexity of 1.25 indicates that the model is moderately confident in its prediction.

5.1.2 Accuracy

Accuracy measures the system's ability to generate correct responses. Mathematically, accuracy is represented as:

$$A = \frac{(TP + TN)}{(TP + TN + FP + FN)}$$

where A is the accuracy, TP is the true positives, TN is the true negatives, FP is the false positives, and FN is the false negatives.

5.1.3 F1-Score

F1-Score measures the system's ability to balance precision and recall. Mathematically, F1-score is represented as:

$$F1 = 2 * \frac{(PR)}{(P + R)}$$

where F1 is the F1-score, P is the precision, and R is the recall.

5.1.4 BLEU Score

BLEU (Bilingual Evaluation Understudy) measures the system's ability to generate fluent and coherent responses. Mathematically, BLEU score is represented as:

$$BLEU = B * \exp\left(\sum \frac{1}{n} \log(P_n)\right)$$

where BLEU is the BLEU score, B is the brevity penalty, n is the n-gram order, and P_n is the precision for each n-gram.

5.1.5 ROUGE Score

ROUGE measures the system's ability to generate responses that align with user expectations. Mathematically, ROUGE score is represented as:

$$ROUGE = \sum \frac{(R * P)}{(R + P)}$$

where ROUGE is the ROUGE score, R is the recall, and P is the precision.

5.1.6 METEOR Score

METEOR (Metric for Evaluation of Translation with Explicit ORdering) score measures the similarity between the predicted response and the reference response. Mathematically, METEOR score is represented as:

$$METEOR = \frac{(Precision * Recall)}{(0.5 * Precision + 0.5 * Recall)}$$

where Precision and Recall are the precision and recall of the predicted response. where ROUGE is the ROUGE score, R is the recall, and P is the precision.

Example: Suppose we have a CAI model that predicts a response with a precision of 0.9 and a recall of 0.8. The METEOR score would be:

$$METEOR = \frac{(0.9 * 0.8)}{(0.5 * 0.9 + 0.5 * 0.8)} \approx 0.88$$

Interpretation: A METEOR score of 0.88 indicates a moderate similarity between the predicted response and the reference response.

5.2 Human Evaluation Metrics

5.2.1 Response Quality

Response Quality measures the system's ability to generate relevant and accurate responses. Mathematically, response quality is represented as:

$$Q = \sum (r_i * w_i)$$

where Q is the response quality, r_i is the response score, and w_i is the weight assigned to each response.

5.2.2 Conversational Flow

Conversational Flow measures the system's ability to engage in natural-sounding conversations. Mathematically, conversational flow is represented as:

$$CF = \sum (c_i * w_i)$$

where CF is the conversational flow, c_i is the conversational flow score, and w_i is the weight assigned to each score.

5.2.3 User Satisfaction

User Satisfaction measures the user's overall satisfaction with the conversation. Mathematically, user satisfaction is represented as:

$$US = \sum (s_i * w_i)$$

where US is the user satisfaction, s_i is the satisfaction score, and w_i is the weight assigned to each score.

5.2.4 Engagement

Engagement measures the user's level of engagement with the conversation. Mathematically, engagement is represented as:

$$E = \sum (e_i * w_i)$$

where E is the engagement, e_i is the engagement score, and w_i is the weight assigned to each score.

6 Applications

Conversational AI has a wide range of applications across various industries, transforming the way businesses operate, interact with customers, and provide services.

6.1 Customer Service

- **Chatbots and Virtual Assistants:** Provide 24/7 customer support, helping customers with queries, orders, and complaints.
- **Automated Issue Resolution:** Use machine learning to identify and resolve common issues, reducing the need for human intervention.
- **Personalized Product Recommendations:** Analyze customer data to suggest relevant products, increasing sales and customer satisfaction.

6.2 Healthcare

- **Virtual Nursing Assistants:** Support patients with medication reminders, appointment scheduling, and basic care instructions.
- **Symptom Checking and Diagnosis:** Use natural language processing to analyze patient symptoms and provide potential diagnoses.
- **Personalized Medication Reminders:** Send patients personalized reminders to take their medication, improving adherence rates.

6.3 Education

- **Virtual Teaching Assistants:** Support teachers with grading, feedback, and student engagement.
- **Adaptive Learning Platforms:** Use machine learning to adjust the difficulty level of course materials based on student performance.
- **Automated Grading and Feedback Systems:** Reduce teacher workload and provide instant feedback to students.

6.4 Finance

- **Chatbots for Account Management:** Help customers with account queries, transactions, and balance inquiries.
- **Virtual Assistants for Investment Advice:** Provide personalized investment advice and portfolio management.

- **Fraud Detection and Prevention Systems:** Use machine learning to detect and prevent fraudulent transactions.

6.5 Entertainment

- **Virtual Assistants for Content Recommendations:** Suggest personalized content, such as movies, TV shows, or music.
- **Chatbots for Interactive Storytelling:** Engage users in interactive stories, such as games or virtual reality experiences.
- **Social Media Platforms:** Use conversational AI to enhance social interaction and community building.

6.6 Transportation

- **Virtual Assistants for Route Planning:** Provide personalized route suggestions based on traffic, time, and location.
- **Chatbots for Ride-Hailing and Logistics:** Help customers with ride-hailing, tracking, and logistics management.
- **Autonomous Vehicles:** Use conversational AI to enhance safety and efficiency in self-driving cars.

6.7 Retail

- **Virtual Assistants for Product Recommendations:** Suggest personalized products based on customer data and preferences.
- **Chatbots for Customer Support:** Provide 24/7 customer support for retail customers.
- **Augmented Reality Shopping Experiences:** Enhance the shopping experience with AR-powered product demonstrations and try-ons.

6.8 Example - Google Assistant

Google Assistant is a leading conversational AI platform that uses natural language processing (NLP) and machine learning (ML) to understand and respond to user requests. With its advanced NLU capabilities, Google Assistant can engage in contextual conversations, understand follow-up questions, and even initiate conversations. Its seamless integration with various Google services, such as Calendar, Maps, and Search, enables users to access a wide range of information and perform tasks hands-free. Additionally, Google Assistant's

compatibility with numerous devices, including Android smartphones, Google Home smart speakers, and iOS devices, makes it a widely accessible conversational AI solution. Overall, Google Assistant's robust capabilities and extensive ecosystem make it a top player in the conversational AI space.

7 Challenges

Despite the advancements in Conversational AI, several challenges remain to be addressed:

7.1 Understanding Natural Language

- **Ambiguity and Context:** Natural language is inherently ambiguous, making it challenging for AI systems to understand the context and intent behind user input.
- **Idioms, Colloquialisms, and Sarcasm:** AI systems struggle to understand idioms, colloquialisms, and sarcasm, which can lead to misinterpretation and incorrect responses.

7.2 Generating Human-Like Responses

- **Lack of Common Sense:** AI systems often lack common sense and real-world experience, making it difficult to generate responses that are relevant and engaging.
- **Emotional Intelligence:** AI systems struggle to understand and replicate human emotions, leading to responses that may come across as insensitive or robotic.

7.3 Handling Multi-Turn Conversations

- **Contextual Understanding:** AI systems need to understand the context of the conversation and remember previous interactions to provide relevant responses.
- **Conversational Flow:** AI systems must manage the conversational flow, knowing when to ask follow-up questions or provide additional information.

7.4 Ensuring Data Quality and Security

- **Data Bias:** AI systems can perpetuate biases present in the training data, leading to unfair and discriminatory responses.
- **Data Security:** Conversational AI systems handle sensitive user data, making data security and protection a top priority.

7.5 Evaluating and Improving Performance

- **Evaluation Metrics:** Developing effective evaluation metrics to measure the performance of Conversational AI systems is an ongoing challenge.
- **Continuous Improvement:** AI systems require continuous training and improvement to stay up-to-date with changing user needs and preferences.

7.6 Ethical Concerns and Potential Biases

- One of the significant challenges and limitations of conversational AI is the potential for biases and ethical concerns. As conversational AI systems learn from vast amounts of data, they can inherit and amplify existing biases, perpetuating discrimination and inequality. For instance, language models may learn to associate certain words or phrases with specific demographics, leading to biased responses.
- Furthermore, conversational AI systems can be designed with a particular worldview or cultural perspective, neglecting diverse voices and experiences.
- Additionally, the lack of transparency and accountability in AI decision-making processes can lead to manipulation and exploitation of users. Addressing these biases and ethical concerns is crucial to ensure that conversational AI systems are fair, inclusive, and respectful of all users.

8 Future Scope

To overcome the challenges in Conversational AI, researchers and developers are exploring new techniques and approaches, including:

8.1 Addressing Current Challenges

- **Multitask Learning:** Training AI systems on multiple tasks simultaneously to improve their understanding of natural language and ability to generate human-like responses.
- **Adversarial Training:** Training AI systems to generate responses that are robust against adversarial attacks and data noise.
- **Human-in-the-Loop:** Involving humans in the training and evaluation process to ensure

AI systems are transparent, explainable, and fair.

- **Ethical Concerns and Biases:** Future research directions should prioritize the development of more transparent, explainable, and fair conversational AI systems. This can be achieved by designing and implementing algorithms that detect and mitigate biases, as well as incorporating diverse perspectives and voices into the development process. Additionally, researchers should explore the creation of more robust and reliable evaluation metrics to assess the ethical implications of conversational AI systems.

8.2 Emerging Trends

- **Explainable AI:** Developing AI systems that provide transparent and interpretable explanations for their responses and decisions.
- **Edge AI:** Deploying AI systems on edge devices, such as smartphones and smart home devices, to reduce latency and improve performance.
- **Multimodal Interaction:** Enabling AI systems to interact with users through multiple modalities, such as text, speech, and gesture.

9 Conclusion

Conversational Artificial Intelligence (CAI) has revolutionized the way humans interact with technology. From chatbots to virtual assistants, CAI has enabled businesses to provide 24/7 customer support, improve user experiences, and increase efficiency. As a rapidly evolving field, CAI continues to advance with breakthroughs in natural language processing, machine learning, and dialogue management, making it increasingly sophisticated. CAI has numerous applications across various industries, including customer service, healthcare, education, and finance, with the potential to transform these sectors and enhance user experiences. However, successful implementation requires careful design and consideration of user needs and goals to ensure effective and engaging conversations. Moreover, the increasing prevalence of CAI raises important ethical and governance considerations, necessitating guidelines and regulations for responsible development and use. Looking ahead, CAI adoption is expected to rise,

with more businesses and organizations leveraging it to improve customer experiences and boost efficiency. Conversational capabilities will become even more advanced, facilitating more natural and effective interactions. As CAI becomes ubiquitous, there will be a growing emphasis on ethics and governance to ensure responsible development and use. In conclusion, Conversational Artificial Intelligence holds immense potential to revolutionize human-technology interactions. As the field continues to evolve, prioritizing user needs, ethical considerations, and responsible development practices is crucial to ensure that CAI benefits society as a whole. By embracing these principles, we can unlock the full potential of CAI and create a brighter future for all.

References

- Raphael Collobert, Jason Weston, Leon Bottou, Michael Karlen, and Koray Kavukcuoglu. 2011. *Natural Language Processing (almost) from Scratch*. Proceedings of the 49th Annual Meeting of the Association for Computational Linguistics (pp. 232-236).
- Ian Goodfellow, Yoshua Bengio, and Aaron Courville. 2016. *Deep Learning* MIT Press
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. *Attention Is All You Need*. Proceedings of the 31st International Conference on Neural Information Processing Systems (pp. 5998-6006)
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, Kristina Toutanova, and Ye Zhang. 2019. *BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding* Proceedings of the 33rd International Conference on Neural Information Processing Systems (pp. 4371-4381).
- Ken Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. *BLEU: A Method for Automatic Evaluation of Machine Translation*. In Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics (pp. 311-318).
- Ken Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. *BLEU: A Method for Automatic Evaluation of Machine Translation*. In Proceedings of the 40th Annual Meeting of the Association for Computational Linguistics (pp. 311-318).